

CS 224d: Assignment #1

Due date: 4/14 11:59 PM PST (You are allowed to use three (3) late days maximum for this assignment)

Note: This is the complementary problem set of Assignment 1 for CS 224d. We strongly recommend that you finish these problems before writing the code for the assignment.

These questions require thought, but do not require long answers. Please be as concise as possible.

We encourage students to discuss in groups for assignments. However, each student finish the problem set and programming assignment individually, and must turn in her/his assignment. We ask that you abide the university Honor Code and that of the Computer Science department, and make sure that all of your submitted work are done by yourself.

Please review any additional instructions posted on the assignment page at <http://cs224d.stanford.edu/assignments.html>. When you are ready to submit, please follow the instructions on the course website.

1 Softmax

Prove that softmax is invariant to constant offsets in the input, that is, for any input vector \mathbf{x} and any constant c ,

$$\text{softmax}(\mathbf{x}) = \text{softmax}(\mathbf{x} + c)$$

where $\mathbf{x} + c$ means adding the constant c to every dimension of \mathbf{x} .

Note: In practice, we make use of this property and choose $c = -\max_i x_i$ when computing softmax probabilities for numerical stability (i.e. subtracting its maximum element from all elements of \mathbf{x}).

2 Neural Network Basics

- (a) Derive the gradients of the sigmoid function and show that it can be rewritten as a function of the function value (i.e. in some expression where only $\sigma(x)$, but not x , is present).
- (b) Derive the gradient of a softmax function when the cost function is cross entropy, i.e. find the gradients with respect to the softmax input vector $\boldsymbol{\theta}$, when the prediction is made by $\hat{\mathbf{y}} = \text{softmax}(\boldsymbol{\theta})$. Remember the cross entropy function is

$$CE(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_i y_i \log(\hat{y}_i)$$


where \mathbf{y} is the one-hot label vector, and $\hat{\mathbf{y}}$ is the predicted probability vector for all classes.

- (c) Derive the gradients with respect to the *inputs* \mathbf{x} to an one-hidden-layer neural network. The neural network employs sigmoid activation function for the hidden layer, and softmax for the output layer. Assume the one-hot label vector is \mathbf{y} , and cross entropy cost is used. (feel free to use $\sigma'(x)$ as the shorthand for sigmoid gradient, and feel free to define any variables whenever you see fit)

Recall that the forward propagation is as follows

$$\mathbf{h} = \text{sigmoid}(\mathbf{x}\mathbf{W}_1 + \mathbf{b}_1)$$

$$\hat{\mathbf{y}} = \text{softmax}(\mathbf{h}\mathbf{W}_2 + \mathbf{b}_2)$$



Note that here we're assuming that the input vector (thus the hidden variables and output probabilities) is a row vector to be consistent with the programming assignment. \mathbf{W}_i and \mathbf{b}_i ($i = 1, 2$) are the weights and biases, respectively, of the two layers.

- (d) How many parameters are there in this neural network, assuming the input is D_x -dimensional, the output is D_y -dimensional, and there are H hidden units?

3 word2vec

- (a) Assume you are given a predicted word vector $\hat{\mathbf{r}}$, and word prediction is made with the softmax function found in word2vec models

$$\Pr(\text{word}_i \mid \hat{\mathbf{r}}, \mathbf{w}) = \frac{\exp(\hat{\mathbf{r}}^\top \mathbf{w}_i)}{\sum_{i'} \exp(\hat{\mathbf{r}}^\top \mathbf{w}_{i'})}$$

where \mathbf{w}_i ($i = 1, \dots, |V|$) are the word vectors for all words in the vocabulary. Assume cross entropy cost is applied to this prediction, derive the gradients with respect to $\hat{\mathbf{r}}$.

- (b) In the previous problem, derive gradients for the word vectors \mathbf{w}_i 's.
- (c) Repeat part (a) and (b) assuming we are using the negative sampling loss for the predicted vector $\hat{\mathbf{r}}$. Assume that K negative samples are drawn, and they are $1, \dots, K$, respectively for simplicity of notation. After you've done this, describe with one sentence why this cost function is much more efficient than the softmax loss (you could provide a speed-up ratio, i.e. the runtime of the softmax loss divided by the runtime of the negative sampling loss).
- (d) Derive gradients for all of the word vectors for skip-gram and CBOW given part (a) and (b), given a set of context words $[\text{word}_{i-C}, \dots, \text{word}_{i-1}, \text{word}_i, \text{word}_{i+1}, \dots, \text{word}_{i+C}]$, where C is the context size. You can denote the word vector for word_k as \mathbf{w}_k for convenience. (*Hint: try to reuse your answers for parts (a) and (b)!*)

4 Sentiment Analysis

- (a) Explain in less than three sentences why do we want to introduce regularization when doing classification (in fact, most machine learning tasks).
- (b) Plot the classification accuracy on the dev set with respect to the regularization value, using a logarithmic scale on the x-axis. Briefly explain with less than three sentences what you see in the plot.