

# **CS224d**

## **Deep NLP**

### **Lecture 4:**

# **Word Window Classification and Neural Networks**

**Richard Socher**

# Overview Today:

- General classification background
- Updating word vectors for classification
- Window classification & cross entropy error derivation tips
- A single layer neural network!
- (Max-Margin loss and **backprop**)

# Refresher: Classification setup and notation

- Generally we have a training dataset consisting of samples


$$\{x_i, y_i\}_{i=1}^N$$

- $x_i$  - inputs, e.g. words (indices or vectors!), context windows, sentences, documents, etc.
- $y_i$  - labels we try to predict, e.g. sentiment, other words, named entities (loc., org. per.), buy/sell decision, later: multi-word sequences

# Classification intuition

- Training data:  $\{x_i, y_i\}_{i=1}^N$
- Simple illustration case:
  - Fixed 2d word vectors to classify
  - Using logistic regression
  - $\rightarrow$  linear decision boundary  $\rightarrow$
- General ML: assume  $x$  is fixed and only train logistic regression weights  $W$  and only modify the decision boundary

$$p(y|x) = \frac{\exp(W_y \cdot x)}{\sum_{c=1}^C \exp(W_c \cdot x)}$$



Visualizations with ConvNetJS by Karpathy  
<http://cs.stanford.edu/people/karpathy/convnetjs/demo/classify2d.html>

# Classification notation


- General ML: only train logistic regression weights and hence only modify the decision boundary
- Loss function over dataset  $\{x_i, y_i\}_{i=1}^N$

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N -\log \left( \frac{e^{f_{y_i}}}{\sum_{c=1}^C e^{f_c}} \right)$$

- Where for each data pair  $(x_i, y_i)$ :
- We can write  $f$  in matrix notation and index elements of it based on class:

$$f_y = f_y(x) = W_y \cdot x = \sum_{j=1}^d W_{yj} x_j$$

$$f = Wx$$




# Classification: Regularization!

- Really full loss function over any dataset includes **regularization** over all parameters  $\theta$ :

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N -\log \left( \frac{e^{f_{y_i}}}{\sum_{c=1}^C e^{f_c}} \right) + \lambda \sum_k \theta_k^2$$

- Regularization will prevent overfitting when we have a lot of features (or later a very powerful/deep model)
  - x-axis: more powerful model or more training iterations
  - Blue: training error, red: test error




# Classification difference with word vectors

- For general machine learning  $\theta$  usually only consists of columns of  $W$ :

$$\theta = \begin{bmatrix} W_{.1} \\ \vdots \\ W_{.d} \end{bmatrix} = W(:,) \in \mathbb{R}^{Cd}$$

- So we only update the decision boundary

$$\nabla_{\theta} J(\theta) = \begin{bmatrix} \nabla_{W_{.1}} \\ \vdots \\ \nabla_{W_{.d}} \end{bmatrix} \in \mathbb{R}^{Cd}$$



Visualizations with ConvNetJS by Karpathy

# Classification difference with word vectors

- For general ML  $\theta$  usually only consists of columns of  $W$
- Additionally common in deep learning:
  - Learn both  $W$  and word vectors  $x$

$$\nabla_{\theta} J(\theta) = \begin{bmatrix} \nabla_{W_{.1}} \\ \vdots \\ \nabla_{W_{.d}} \\ \nabla_{x_{aardvark}} \\ \vdots \\ \nabla_{x_{zebra}} \end{bmatrix} \in \mathbb{R}^{Cd + Vd}$$


Very large!

Overfitting Danger!




# Loosing generalization by re-training word vectors

- Setting: Training logistic regression for movie review sentiment and in the training data we have the word
  - “TV” and “telly”
- In the testing data we have
  - “television”
- Originally they were all similar (from pre-training word vectors)
- What happens when we train the word vectors?



# Loosing generalization by re-training word vectors

- What happens when we train the word vectors?
  - Those that are in the training data move around
  - Words from pre-training that do NOT appear in training stay
- Example:
- In training data: “TV” and “telly”
- In testing data only: “television”




# Loosing generalization by re-training word vectors

- Take home message:

If you only have a small training data set, don't train the word vectors.

If you have have a very large dataset, it may work better to train word vectors to the task.



## Side note on word vectors notation

- The word vector matrix  $L$  is also called lookup table
- Word vectors = word embeddings = word representations (mostly)
- Mostly from methods like word2vec or Glove

$$L = d \begin{bmatrix} \text{aardvark} & \text{a} & \dots & \text{meta} & \dots & \text{zebra} \end{bmatrix}^{|V|}$$

- These are the word features  $x_{\text{word}}$  from now on
- Conceptually you get a word's vector by left multiplying a one-hot vector  $e$  by  $L$ :  $x = Le \in d \times V \cdot V \times 1$

# Window classification

- Classifying single words is rarely done.
- Interesting problems like ambiguity arise in context!
- Example: auto-antonyms:
  - "To sanction" can mean "to permit" or "to punish."
  - "To seed" can mean "to place seeds" or "to remove seeds."
- Example: ambiguous named entities:
  - Paris → Paris Hilton vs Paris, France
  - Hathaway → Berkshire Hathaway, Anne Hathaway

# Window classification

- Idea: Instead of classifying a single word, just classify a word together with its context window of neighboring words.
- For example named entity recognition into 4 classes:
  - Person, location, organization, none
- Many possibilities exist for classifying one word in context, e.g. averaging all the words in a window but that loses position information

# Window classification

- Most commonly used technique to classify a word in a window
- Train classifier by assigning a label to a center word and concatenating all word vectors surrounding it.
- Example: Classify Paris in the context of this sentence with window length 2:

... museums in Paris are amazing ...

● ● ● ●   ● ● ● ●   ● ● ● ●   ● ● ● ●   ● ● ● ●

$$X_{\text{window}} = \begin{bmatrix} x_{\text{museums}} & x_{\text{in}} & x_{\text{Paris}} & x_{\text{are}} & x_{\text{amazing}} \end{bmatrix}$$

- Resulting vector  $x_{\text{window}} = \boxed{x \in \mathbb{R}^{5d}}$ , a column vector!

# Simplest window classifier: Softmax

- With  $x = x_{\text{window}}$  we can use the same softmax classifier as before

= predicted model  
output probability

$$\hat{y} = p(y|x) = \frac{\exp(W_y \cdot x)}{\sum_{c=1}^C \exp(W_c \cdot x)}$$

same

- With cross entropy error as before:

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N -\log \left( \frac{e^{f_{y_i}}}{\sum_{c=1}^C e^{f_c}} \right)$$

- But how do you update the word vectors?



# Updating concatenated word vectors

- Short answer: Just take derivatives as before
- Long answer: Let's go over the steps together (you'll have to fill in the details in PSet 1!)
- Define:
  - $\hat{y}$  : softmax probability output vector (see previous slide)
  - $t$  : target probability distribution (all 0's except at ground truth index of class  $y$ , where it's 1)
  - $f = Wx \in \mathbb{R}^C$  and  $f_c = c$ 'th element of the  $f$  vector
- Hard, the first time, hence some tips now :)

# Updating concatenated word vectors

- Tip 1: Carefully define your variables and keep track of their dimensionality!  $f = f(x) = Wx \in \mathbb{R}^C$   
 $\hat{y} \quad t \quad W \in \mathbb{R}^{C \times 5d}$
- Tip 2: **Know thy chain rule** and don't forget in which variables other variables are being used:

$$\frac{\partial}{\partial x} - \log \text{softmax}(f_y(x)) = \sum_{c=1}^C - \frac{\partial \log \text{softmax}(f_y(x))}{\partial f_c} \cdot \frac{\partial f_c(x)}{\partial x}$$

- Tip 3: For the softmax part of the derivative: First take the derivative wrt  $f_c$  when  $c=y$  (the correct class), then take derivative wrt  $f_c$  when  $c \neq y$  (all the incorrect classes)

# Updating concatenated word vectors

- Tip 4: When you take derivative wrt one element of  $f$ , try to see if you can create a gradient in the end that includes all partial derivatives:

$$\hat{y} \quad t$$
$$f = f(x) = Wx \in \mathbb{R}^C$$

$$\frac{\partial}{\partial f} - \log \text{softmax}(f_y) = \begin{bmatrix} \hat{y}_1 \\ \vdots \\ \hat{y}_y - 1 \\ \vdots \\ \hat{y}_C \end{bmatrix}$$

- Tip 5: To later not go insane, think of your results in terms of vector operations and define new, single index-able vectors:

$$\frac{\partial}{\partial f} - \log \text{softmax}(f_y) = [\hat{y} - t] = \delta$$

# Updating concatenated word vectors

- Tip 5: When you start with the chain rule, first use explicit sums and look at partial derivatives of e.g.  $x_i$  or  $W_{ij}$

$$\hat{y} = t$$
$$f = f(x) = Wx \in \mathbb{R}^C$$

$$\sum_{c=1}^C -\frac{\partial \log \text{softmax}(f_y(x))}{\partial f_c} \cdot \frac{\partial f_c(x)}{\partial x} = \sum_{c=1}^C \delta_c W_c.$$

- Tip 6: To clean it up for even more complex functions later: Know dimensionality of variables & simplify into matrix notation

$$\frac{\partial}{\partial x} -\log p(y|x) = \sum_{c=1}^C \delta_c W_c. = W^T \delta$$

- Tip 7: Write this out in full sums if it's not clear!

# Updating concatenated word vectors

- Tip 5: When you start with the chain rule, first use explicit sums and look at partial derivatives of e.g.  $x_i$  or  $W_{ij}$

$$\hat{y} = t$$
$$f = f(x) = Wx \in \mathbb{R}^C$$

$$\sum_{c=1}^C -\frac{\partial \log \text{softmax}(f_y(x))}{\partial f_c} \cdot \frac{\partial f_c(x)}{\partial x} = \sum_{c=1}^C \delta_c W_c.$$

- Tip 6: To clean it up for even more complex functions later: Know dimensionality of variables & simplify into matrix notation

$$\frac{\partial}{\partial x} - \log p(y|x) = \sum_{c=1}^C \delta_c W_c = W^T \delta$$

- Tip 7: Write this out in full sums if it's not clear!

# Updating concatenated word vectors

- What is the dimensionality of

$$\frac{\partial}{\partial x} - \log p(y|x) = \sum_{c=1}^C \delta_c W_c. = W^T \delta$$

- $X$  is the entire window of 5 d-dimensional word vectors, so the derivative wrt to  $x$  has to have the same dimensionality:

$$\nabla_x J = W^T \delta \in \mathbb{R}^{5d}$$

# Updating concatenated word vectors

- The gradient that arrives at and updates the word vectors can simply be split up for each word vector:
- Let  $\nabla_x J = W^T \delta = \delta_{window}$
- With  $x_{window} = [x_{museums} \quad x_{in} \quad x_{Paris} \quad x_{are} \quad x_{amazing}]$

- We have

$$\delta_{window} = \begin{bmatrix} \nabla_{x_{museums}} \\ \nabla_{x_{in}} \\ \nabla_{x_{Paris}} \\ \nabla_{x_{are}} \\ \nabla_{x_{amazing}} \end{bmatrix} \in \mathbb{R}^{5d}$$

# Updating concatenated word vectors

- This will push word vectors into areas such they will be helpful in determining named entities.
- For example, the model can learn that seeing  $x_{in}$  as the word just before the center word is indicative for the center word to be a location



# What's missing for training the window model?

- The gradient of  $J$  wrt the softmax weights  $W$ !
- Similar steps, write down partial wrt  $W_{ij}$  first!
- Then we have full

$$\nabla_{\theta} J(\theta) = \begin{bmatrix} \nabla_{W_{\cdot 1}} \\ \vdots \\ \nabla_{W_{\cdot d}} \\ \nabla_{x_{aardvark}} \\ \vdots \\ \nabla_{x_{zebra}} \end{bmatrix} \in \mathbb{R}^{Cd+Vd}$$

# A note on matrix implementations

- There are two expensive operations in the softmax:
- The matrix multiplication  $f = Wx$  and the exp
- A for loop is never as efficient when you implement it compared vs when you use a larger matrix multiplication that does the same mathematical operation!
- Example code →

# A note on matrix implementations

- Looping over word vectors instead of concatenating them all into one large matrix and then multiplying the softmax weights with that matrix

```
from numpy import random
N = 500 # number of windows to classify
d = 300 # dimensionality of each window
C = 5 # number of classes
W = random.rand(C,d)
wordvectors_list = [random.rand(d,1) for i in range(N)]
wordvectors_one_matrix = random.rand(d,N)

%timeit [W.dot(wordvectors_list[i]) for i in range(N)]
%timeit W.dot(wordvectors_one_matrix)
```

- 1000 loops, best of 3: 639  $\mu$ s per loop  
10000 loops, best of 3: 53.8  $\mu$ s per loop

# A note on matrix implementations

```
from numpy import random
N = 500 # number of windows to classify
d = 300 # dimensionality of each window
C = 5 # number of classes
W = random.rand(C,d)
wordvectors_list = [random.rand(d,1) for i in range(N)]
wordvectors_one_matrix = random.rand(d,N)

%timeit [W.dot(wordvectors_list[i]) for i in range(N)]
%timeit W.dot(wordvectors_one_matrix)
```


- Result of faster method is a  $C \times N$  matrix:
  - Each column is an  $f(x)$  in our notation (unnormalized class scores)
- Matrices are awesome!
- You should speed test your code a lot too

## Softmax (= logistic regression) is not very powerful

- Softmax only gives linear decision boundaries in the original space.
- With little data that can be a good regularizer
- With more data it is very limiting!

# Softmax (= logistic regression) is not very powerful

- Softmax only linear decision boundaries





→ Lame when problem is complex

Wouldn't it be cool to get these correct?

# Neural Nets for the Win!

- Neural networks can learn much more complex functions and nonlinear decision boundaries!



# From logistic regression to neural nets




# Demystifying neural networks

Neural networks come with their own terminological baggage

... just like SVMs

But if you understand how softmax models work

Then **you already understand** the operation of a basic neural network neuron!





# A neuron is essentially a binary logistic regression unit

$$h_{w,b}(x) = f(w^T x + b)$$

$b$ : We can have an “always on” feature, which gives a class prior, or separate it out, as a bias term

$$f(z) = \frac{1}{1 + e^{-z}}$$




$w, b$  are the parameters of this neuron  
i.e., this logistic regression model

# A neural network

= running several logistic regressions at the same time

If we feed a vector of inputs through a bunch of logistic regression functions, then we get a vector of outputs ...




*But we don't have to decide ahead of time what variables these logistic regressions are trying to predict!*

# A neural network

= running several logistic regressions at the same time

... which we can feed into another logistic regression function




*It is the loss function that will direct what the intermediate hidden variables should be, so as to do a good job at predicting the targets for the next layer, etc.*

# A neural network

= running several logistic regressions at the same time

Before we know it, we have a multilayer neural network....



# Matrix notation for a layer

We have

$$a_1 = f(W_{11}x_1 + W_{12}x_2 + W_{13}x_3 + b_1)$$

$$a_2 = f(W_{21}x_1 + W_{22}x_2 + W_{23}x_3 + b_2)$$

etc.


In matrix notation

$$z = Wx + b$$

$$a = f(z)$$



where  $f$  is applied element-wise:

$$f([z_1, z_2, z_3]) = [f(z_1), f(z_2), f(z_3)]$$



# Non-linearities (f): Why they're needed

- Example: function approximation, e.g., regression or classification
  - Without non-linearities, deep neural networks can't do anything more than a linear transform
  - Extra layers could just be compiled down into a single linear transform:
$$W_1 W_2 x = Wx$$
  - With more layers, they can approximate more complex functions!



# A more powerful window classifier

- Revisiting
- $X_{\text{window}} = [x_{\text{museums}} \quad x_{\text{in}} \quad x_{\text{Paris}} \quad x_{\text{are}} \quad x_{\text{amazing}}]$



# A Single Layer Neural Network

- A single layer is a combination of a linear layer and a nonlinearity:

$$z = Wx + b$$

$$a = f(z)$$

- The neural activations  $a$  can then be used to compute some function
- For instance, a softmax probability or an unnormalized score or a we care about:

$$score(x) = U^T a \in \mathbb{R}$$

# Summary: Feed-forward Computation


Computing a window's score with a 3-layer neural net:  $s = \text{score}(\text{museums in Paris are amazing})$

$$s = U^T f(Wx + b) \quad x \in \mathbb{R}^{20 \times 1}, W \in \mathbb{R}^{8 \times 20}, U \in \mathbb{R}^{8 \times 1}$$

$$s = U^T a$$

$$a = f(z)$$

$$z = Wx + b$$



## Next lecture:

Training a window-based neural network.

Taking more **deeper derivatives** → **Backprop**

Then we have all the basic tools in place to learn about more complex models :)