

# CS 224D: DEEP LEARNING FOR NLP

## MIDTERM REVIEW

---

Neural Networks: Terminology, Forward Pass, Backpropagation.

Rohit Mundra

May 4, 2015


# Overview

- Neural Network Example
- Terminology
- Example 1:
  - Forward Pass
  - Backpropagation Using Chain Rule
  - What is delta? From Chain Rule to Modular Error Flow
- Example 2:
  - Forward Pass
  - Backpropagation




# Neural Networks

- One of many different types of non-linear classifiers (i.e. leads to non-linear decision boundaries)




- Most common design involves the stacking of affine transformations followed by point-wise (element-wise) non-linearity

# An example of a neural network



- This is a 4 layer neural network.
- 2 hidden-layer neural network.
- 2-10-10-3 neural network (complete architecture defn.)


# Our first example



- This is a 3 layer neural network
- 1 hidden-layer neural network


# Our first example:

## Terminology





# Our first example:

## Terminology



# Our first example:

## Activation Unit Terminology




$$z_1^{(2)} = W_{11}^{(1)}a_1^{(1)} + W_{12}^{(1)}a_2^{(1)} + W_{13}^{(1)}a_3^{(1)} + W_{14}^{(1)}a_4^{(1)}$$

$a_1^{(2)}$  is the 1<sup>st</sup> activation unit of layer 2

$$a_1^{(2)} = \sigma(z_1^{(2)})$$




# Our first example: Forward Pass



$$\begin{aligned}z_1^{(1)} &= x_1 \\z_2^{(1)} &= x_2 \\z_3^{(1)} &= x_3 \\z_4^{(1)} &= x_4\end{aligned}$$

# Our first example: Forward Pass




$$a_1^{(1)} = z_1^{(1)}$$

$$a_2^{(1)} = z_2^{(1)}$$

$$a_3^{(1)} = z_3^{(1)}$$

$$a_4^{(1)} = z_4^{(1)}$$


# Our first example: Forward Pass



$$z_1^{(2)} = W_{11}^{(1)} a_1^{(1)} + W_{12}^{(1)} a_2^{(1)} + W_{13}^{(1)} a_3^{(1)} + W_{14}^{(1)} a_4^{(1)}$$

$$z_2^{(2)} = W_{21}^{(1)} a_1^{(1)} + W_{22}^{(1)} a_2^{(1)} + W_{23}^{(1)} a_3^{(1)} + W_{24}^{(1)} a_4^{(1)}$$


# Our first example: Forward Pass



$$z^{(2)} = W^{(1)} a^{(1)}$$

$$\begin{bmatrix} z_1^{(2)} \\ z_2^{(2)} \end{bmatrix} = \begin{bmatrix} W_{11}^{(1)} & W_{12}^{(1)} & W_{13}^{(1)} & W_{14}^{(1)} \\ W_{21}^{(1)} & W_{22}^{(1)} & W_{23}^{(1)} & W_{24}^{(1)} \end{bmatrix} \begin{bmatrix} a_1^{(1)} \\ a_2^{(1)} \\ a_3^{(1)} \\ a_4^{(1)} \end{bmatrix}$$


# Our first example: Forward Pass



$$z^{(2)} = W^{(1)} a^{(1)}$$

Affine transformation


# Our first example: Forward Pass



$$a^{(2)} = \boldsymbol{\sigma}(z^{(2)})$$

Point-wise/Element-wise non-linearity


# Our first example: Forward Pass



$$z^{(3)} = W^{(2)} a^{(2)}$$

Affine transformation

# Our first example: Forward Pass




$$a^{(3)} = z^{(3)}$$

$$s = a^{(3)}$$




# Our first example: Backpropagation using chain rule



Let us try to calculate the error gradient wrt  $W_{14}^{(1)}$   
Thus we want to find:

$$\frac{\partial s}{\partial W_{14}^{(1)}}$$


# Our first example: Backpropagation using chain rule



Let us try to calculate the error gradient wrt  $W_{14}^{(1)}$   
Thus we want to find:

$$\frac{\partial s}{\partial z_1^{(3)}} \frac{\partial z_1^{(3)}}{\partial a_1^{(2)}} \frac{\partial a_1^{(2)}}{\partial z_1^{(2)}} \frac{\partial z_1^{(2)}}{\partial W_{14}^{(1)}}$$


# Our first example: Backpropagation using chain rule



This is simply 1

$$\frac{\partial s}{\partial z_1^{(3)}} \frac{\partial z_1^{(3)}}{\partial a_1^{(2)}} \frac{\partial a_1^{(2)}}{\partial z_1^{(2)}} \frac{\partial z_1^{(2)}}{\partial W_{14}^{(1)}}$$


# Our first example: Backpropagation using chain rule



$$\frac{\partial z_1^{(3)}}{\partial a_1^{(2)}} \frac{\partial a_1^{(2)}}{\partial z_1^{(2)}} \frac{\partial z_1^{(2)}}{\partial W_{14}^{(1)}}$$


$$\frac{\partial (W_{11}^{(2)} a_1^{(2)} + W_{12}^{(2)} a_2^{(2)})}{\partial a_1^{(2)}} \frac{\partial a_1^{(2)}}{\partial z_1^{(2)}} \frac{\partial z_1^{(2)}}{\partial W_{14}^{(1)}}$$

# Our first example: Backpropagation using chain rule




$$W_{11}^{(2)} \frac{\partial a_1^{(2)}}{\partial z_1^{(2)}} \frac{\partial z_1^{(2)}}{\partial W_{14}^{(1)}}$$

# Our first example: Backpropagation using chain rule




$$W_{11}^{(2)} \sigma' \left( z_1^{(2)} \right) \frac{\partial z_1^{(2)}}{\partial W_{14}^{(1)}}$$

# Our first example: Backpropagation using chain rule



$$W_{11}^{(2)} \sigma' \left( z_1^{(2)} \right) \frac{\partial (W_{11}^{(1)} a_1^{(1)} + W_{12}^{(1)} a_2^{(1)} + W_{13}^{(1)} a_3^{(1)} + W_{14}^{(1)} a_4^{(1)})}{\partial W_{14}^{(1)}}$$

# Our first example: Backpropagation using chain rule




$$\underbrace{W_{11}^{(2)} \sigma' \left( z_1^{(2)} \right) a_4^{(1)}}_{\delta_1^{(2)}}$$



# Our first example: Backpropagation Observations

We got error  
gradient wrt  $W_{14}^{(1)}$



Required:


- the signal forwarded by  $W_{14}^{(1)} = a_4^{(1)}$
- the error propagating backwards  $W_{11}^{(2)}$
- the local gradient  $\sigma'(z_1^{(2)})$

# Our first example: Backpropagation Observations

We tried to get error  
gradient wrt  $W_{14}^{(1)}$

Required:

- the signal forwarded by  $W_{14}^{(1)} = a_4^{(1)}$
- the error propagating backwards  $W_{11}^{(2)}$
- the local gradient  $\sigma'(z_1^{(2)})$



We can do this for  
all of  $W^{(1)}$ :



$$\begin{pmatrix} \delta_1^{(2)} a_1^{(1)} & \delta_1^{(2)} a_2^{(1)} & \delta_1^{(2)} a_3^{(1)} & \delta_1^{(2)} a_4^{(1)} \\ \delta_2^{(2)} a_1^{(1)} & \delta_2^{(2)} a_2^{(1)} & \delta_2^{(2)} a_3^{(1)} & \delta_2^{(2)} a_4^{(1)} \end{pmatrix}$$

(as outer product)

$$\begin{pmatrix} \delta_1^{(2)} \\ \delta_2^{(2)} \end{pmatrix} \begin{pmatrix} a_1^{(1)} & a_2^{(1)} & a_3^{(1)} & a_4^{(1)} \end{pmatrix}$$

# Our first example:

## Let us define $\delta$



$\delta_1^{(2)}$  is the error flowing backwards at the same point where  $z_1^{(2)}$  passed forwards. Thus it is simply the gradient of the error wrt  $z_1^{(2)}$ .


# Our first example:

## Backpropagation using error vectors



The chain rule of differentiation just boils down very simple patterns in error backpropagation:

1. An error  $x$  flowing backwards passes a neuron by getting amplified by the local gradient.
2. An error  $\delta$  that needs to go through an affine transformation distributes itself in the way signal combined in forward pass.

Orange = Backprop.  
Green = Fwd. Pass

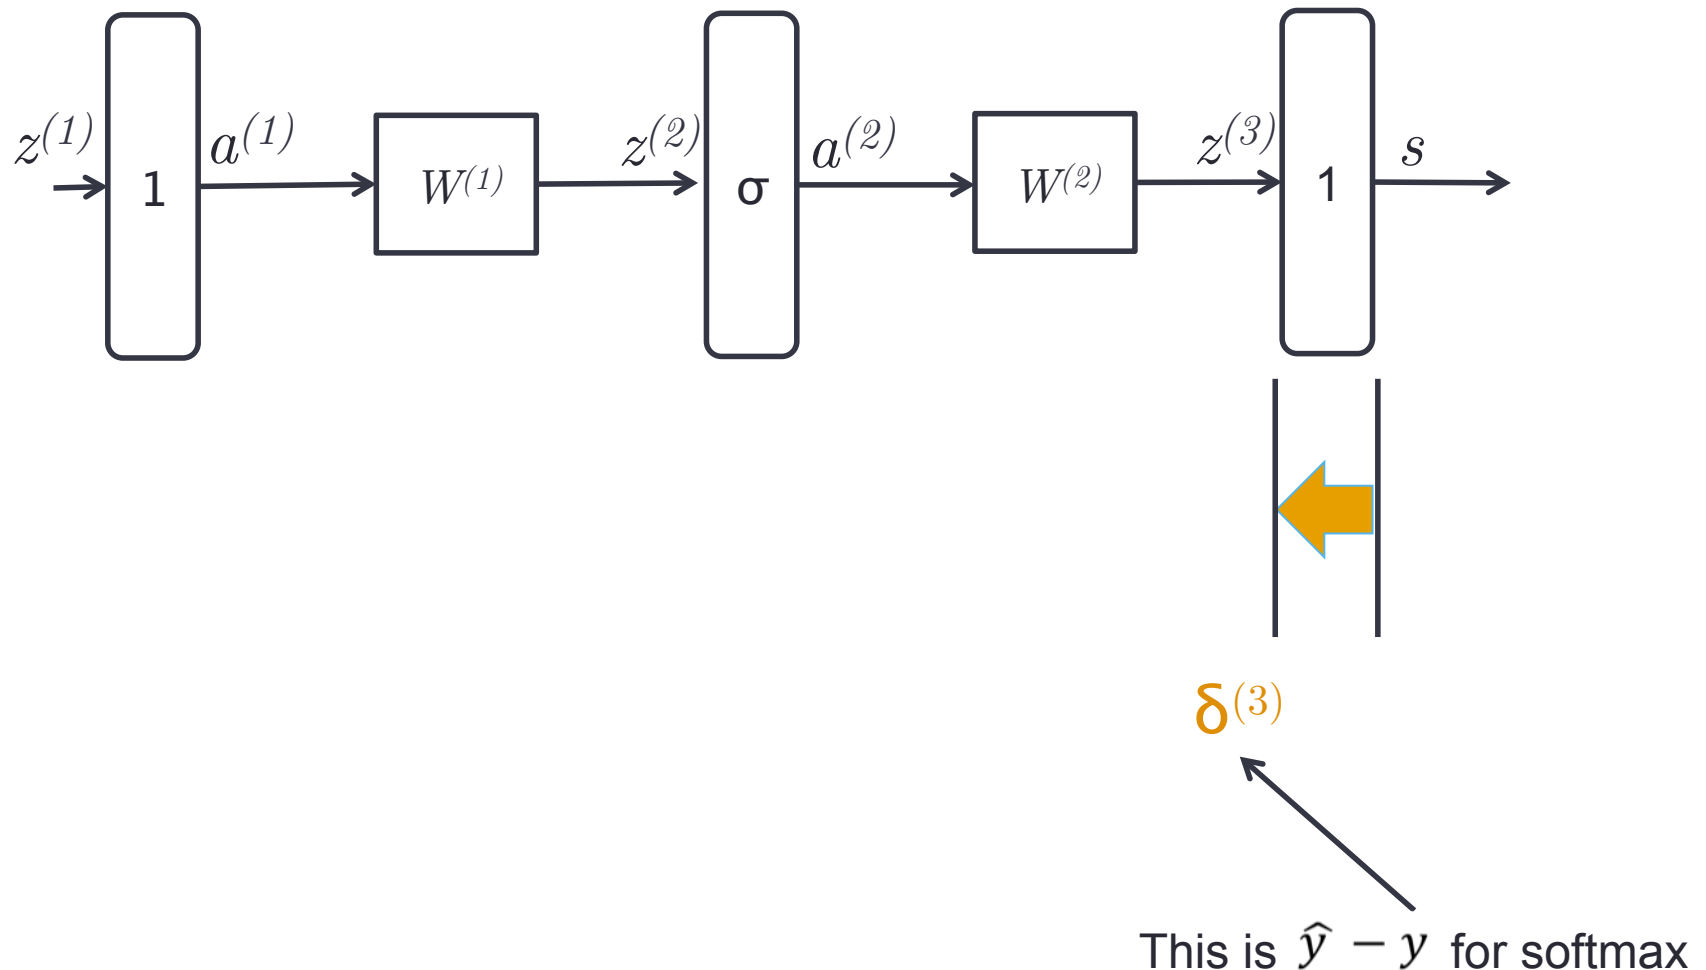


# Our first example: Backpropagation using error vectors




# Our first example:

## Backpropagation using error vectors



# Our first example:


## Backpropagation using error vectors



Gradient w.r.t  $W^{(2)} = \delta^{(3)} a^{(2)T}$

# Our first example:

## Backpropagation using error vectors




- Reusing the  $\delta^{(3)}$  for downstream updates.
- Moving error vector across affine transformation simply requires multiplication with the transpose of forward matrix
- Notice that the dimensions will line up perfectly too!




# Our first example:

## Backpropagation using error vectors



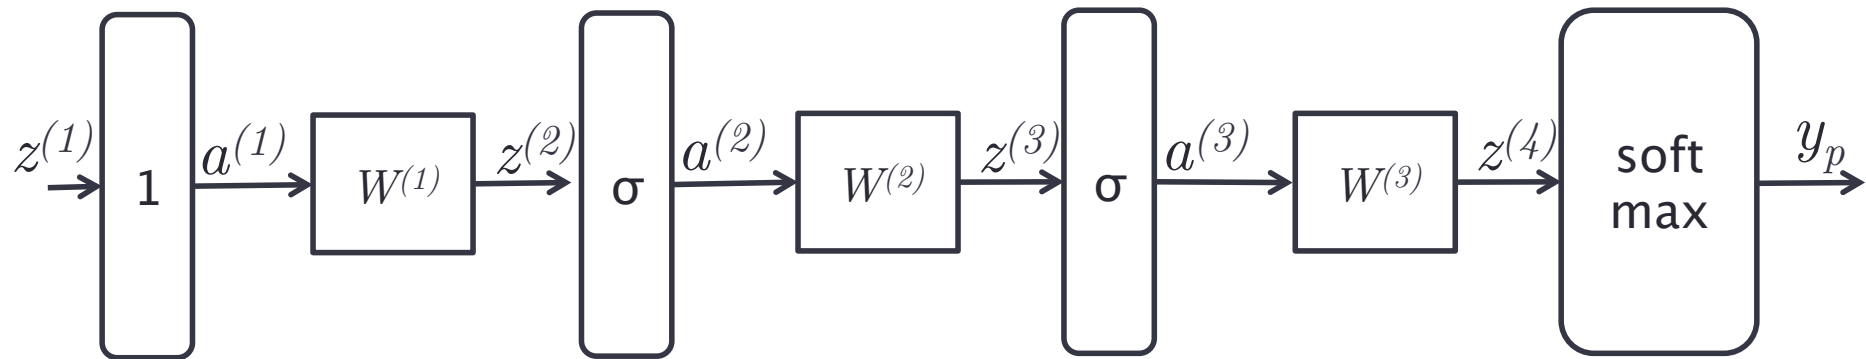
--Moving error vector across point-wise non-linearity requires point-wise multiplication with local gradient of the non-linearity

# Our first example: Backpropagation using error vectors

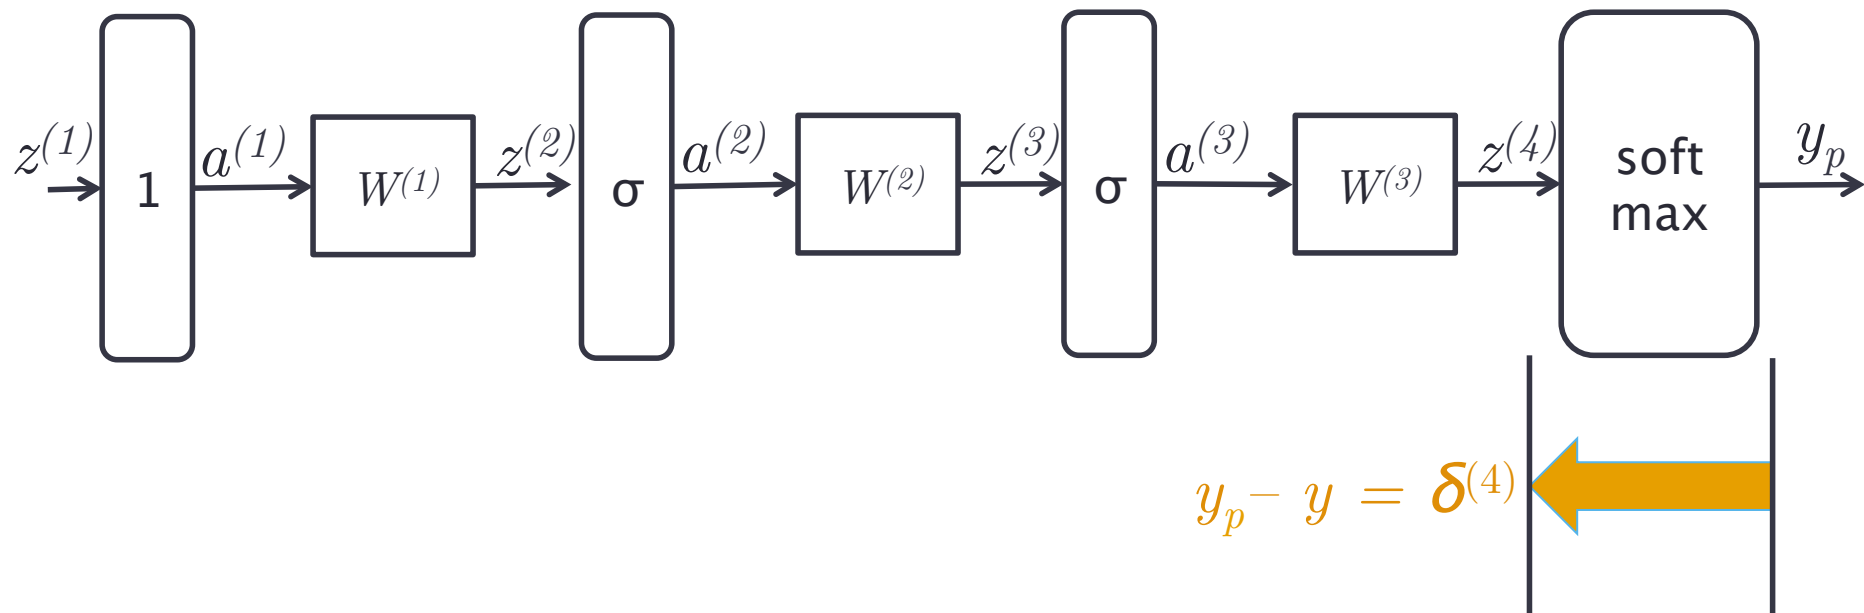


Gradient w.r.t  $W^{(1)} = \delta^{(2)} a^{(1)T}$

# Our second example (4-layer network): Backpropagation using error vectors




# Our second example (4-layer network): Backpropagation using error vectors

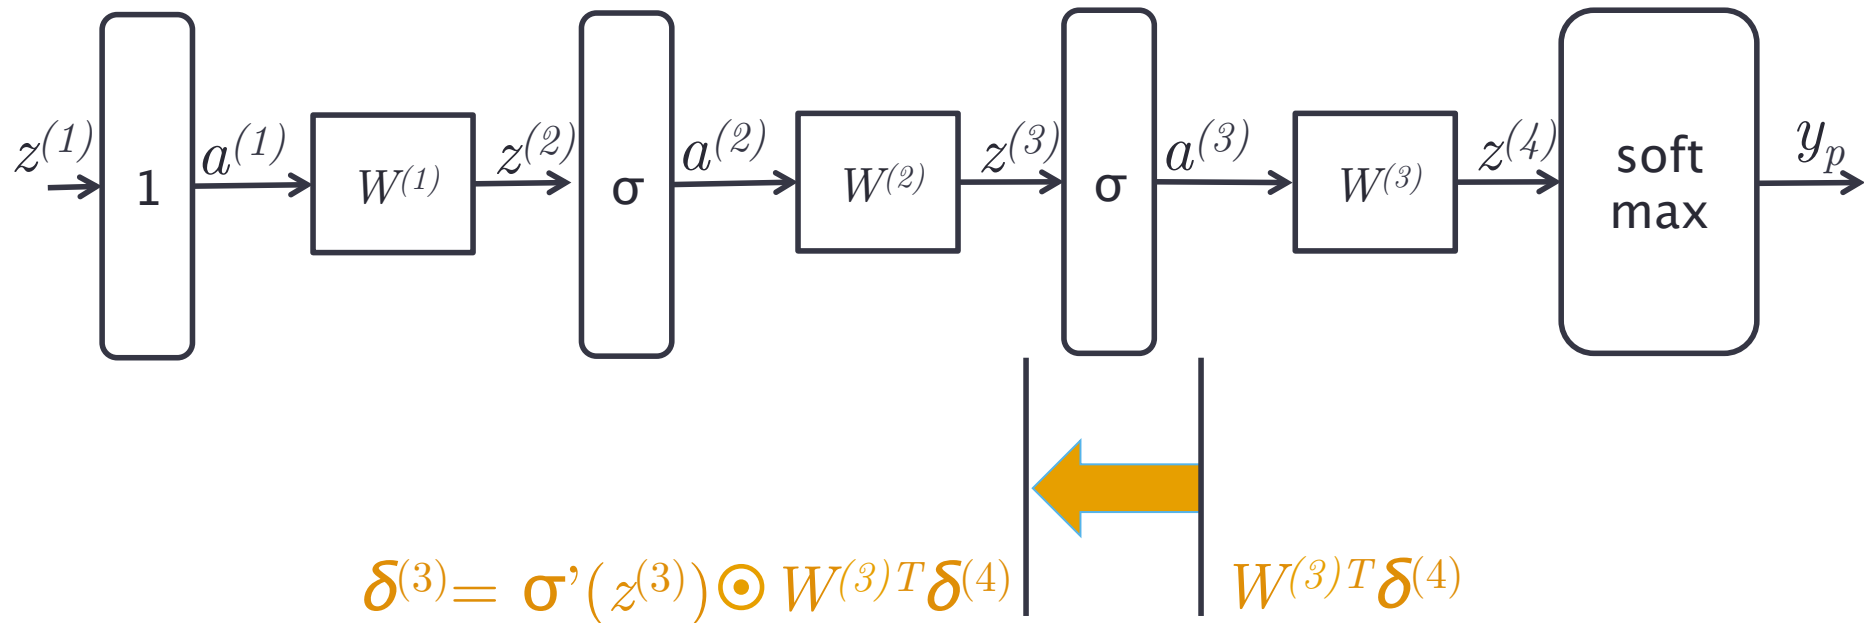


# Our second example (4-layer network): Backpropagation using error vectors

$$\text{Grad } W^{(3)} = \delta^{(4)} a^{(3)T}$$




# Our second example (4-layer network): Backpropagation using error vectors

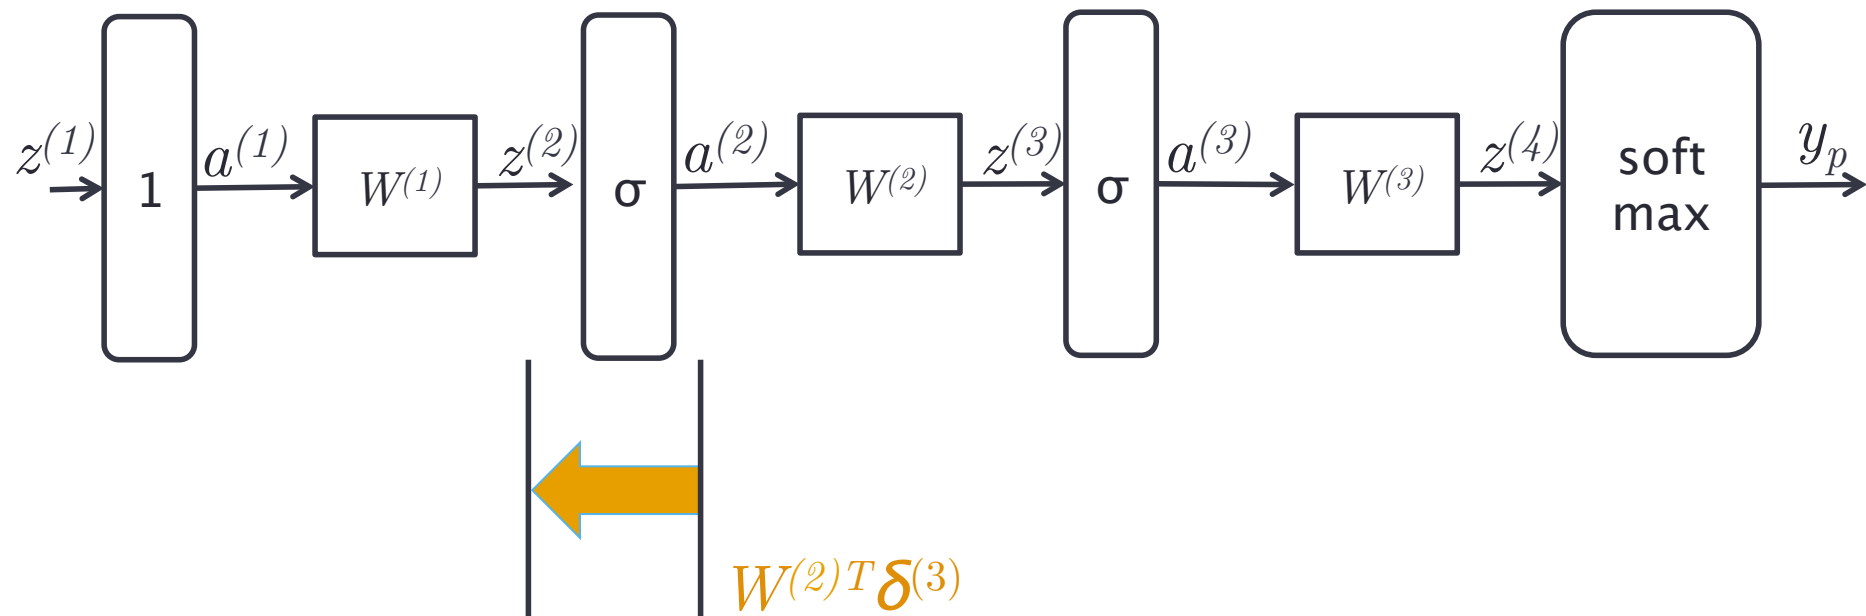


# Our second example (4-layer network): Backpropagation using error vectors

$$\text{Grad } W^{(2)} = \delta^{(3)} a^{(2)T}$$



# Our second example (4-layer network): Backpropagation using error vectors




$$\delta^{(2)} = \sigma'(z^{(2)}) \odot W^{(2)T} \delta^{(3)}$$




# Our second example (4-layer network): Backpropagation using error vectors

$$\text{Grad } W^{(1)} = \delta^{(2)} a^{(1)T}$$



# Our second example (4-layer network): Backpropagation using error vectors

Grad wrt input vector =  $W^{(1)T}\delta^{(2)}$



# CS224D Midterm Review

Ian Tenney

May 4, 2015

# Outline

## Backpropagation (continued)

- RNN Structure

- RNN Backpropagation

## Backprop on a DAG

- Example: Gated Recurrent Units (GRUs)

- GRU Backpropagation

# Outline

## Backpropagation (continued)

- RNN Structure


- RNN Backpropagation

## Backprop on a DAG

- Example: Gated Recurrent Units (GRUs)


- GRU Backpropagation

# Basic RNN Structure



- ▶ Basic RNN ("Elman network")
- ▶ You've seen this on Assignment #2 (and also in Lecture #5)

# Basic RNN Structure




- Two layers between input and prediction, plus hidden state

$$h^{(t)} = \text{sigmoid} \left( Hh^{(t-1)} + Wx^{(t)} + b_1 \right)$$

$$\hat{y}^{(t)} = \text{softmax} \left( Uh^{(t)} + b_2 \right)$$

# Unrolled RNN



- ▶ Helps to think about as “unrolled” network: distinct nodes for each timestep
- ▶ Just do backprop on this! Then combine shared gradients.



# Backprop on RNN

- Usual cross-entropy loss ( $k$ -class):

$$\bar{P}(y^{(t)} = j \mid x^{(t)}, \dots, x^{(1)}) = \hat{y}_j^{(t)}$$
$$J^{(t)}(\theta) = - \sum_{j=1}^k y_j^{(t)} \log \hat{y}_j^{(t)}$$

- Just do backprop on this! First timestep ( $\tau = 1$ ):

$$\frac{\partial J^{(t)}}{\partial U} \quad \frac{\partial J^{(t)}}{\partial b_2}$$
$$\left. \frac{\partial J^{(t)}}{\partial H} \right|_{(t)} \quad \frac{\partial J^{(t)}}{\partial h^{(t)}} \quad \left. \frac{\partial J^{(t)}}{\partial W} \right|_{(t)} \quad \frac{\partial J^{(t)}}{\partial x^{(t)}}$$

# Backprop on RNN

- First timestep ( $s = 0$ ):

$$\frac{\partial J^{(t)}}{\partial U} \quad \frac{\partial J^{(t)}}{\partial b_2}$$
$$\left. \frac{\partial J^{(t)}}{\partial H} \right|_{(t)} \quad \left. \frac{\partial J^{(t)}}{\partial h^{(t)}} \right|_{(t)} \quad \left. \frac{\partial J^{(t)}}{\partial W} \right|_{(t)} \quad \frac{\partial J^{(t)}}{\partial x^{(t)}}$$

- Back in time ( $s = 1, 2, \dots, \tau - 1$ )

$$\left. \frac{\partial J^{(t)}}{\partial H} \right|_{(t-s)} \quad \left. \frac{\partial J^{(t)}}{\partial h^{(t-s)}} \right|_{(t-s)} \quad \left. \frac{\partial J^{(t)}}{\partial W} \right|_{(t-s)} \quad \frac{\partial J^{(t)}}{\partial x^{(t-s)}}$$

Yuck, that's a lot of math!

- ▶ Actually, it's not so bad.
- ▶ Solution: error vectors ( $\delta$ )

# Making sense of the madness

- ▶ Chain rule to the rescue!
- ▶  $a^{(t)} = Uh^{(t)} + b_2$
- ▶  $\hat{y}^{(t)} = \text{softmax}(a^{(t)})$
- ▶ Gradient is *transpose* of Jacobian:

$$\nabla_a J = \left( \frac{\partial J^{(t)}}{\partial a^{(t)}} \right)^T = \hat{y}^{(t)} - y^{(t)} = \delta^{(2)(t)} \in \mathbb{R}^{k \times 1}$$

- ▶ Now dimensions work out:

$$\frac{\partial J^{(t)}}{\partial a^{(t)}} \cdot \frac{\partial a^{(t)}}{\partial b_2} = (\delta^{(2)(t)})^T I \in \mathbb{R}^{(1 \times k) \cdot (k \times k)} = \mathbb{R}^{1 \times k}$$

# Making sense of the madness

- ▶ Chain rule to the rescue!
- ▶  $a^{(t)} = Uh^{(t)} + b_2$
- ▶  $\hat{y}^{(t)} = \text{softmax}(a^{(t)})$
- ▶ Gradient is *transpose* of Jacobian:

$$\nabla_a J = \left( \frac{\partial J^{(t)}}{\partial a^{(t)}} \right)^T = \hat{y}^{(t)} - y^{(t)} = \delta^{(2)}(t) \in \mathbb{R}^{k \times 1}$$

- ▶ Now dimensions work out:

$$\frac{\partial J^{(t)}}{\partial a^{(t)}} \cdot \frac{\partial a^{(t)}}{\partial b_2} = (\delta^{(2)}(t))^T I \in \mathbb{R}^{(1 \times k) \cdot (k \times k)} = \mathbb{R}^{1 \times k}$$

# Making sense of the madness

- ▶ Chain rule to the rescue!
- ▶  $a^{(t)} = Uh^{(t)} + b_2$
- ▶  $\hat{y}^{(t)} = \text{softmax}(a^{(t)})$
- ▶ Gradient is *transpose* of Jacobian:

$$\nabla_a J = \left( \frac{\partial J^{(t)}}{\partial a^{(t)}} \right)^T = \hat{y}^{(t)} - y^{(t)} = \delta^{(2)(t)} \in \mathbb{R}^{k \times 1}$$

- ▶ Now dimensions work out:

$$\frac{\partial J^{(t)}}{\partial a^{(t)}} \cdot \frac{\partial a^{(t)}}{\partial b_2} = (\delta^{(2)(t)})^T I \in \mathbb{R}^{(1 \times k) \cdot (k \times k)} = \mathbb{R}^{1 \times k}$$

# Making sense of the madness

- ▶ Chain rule to the rescue!
- ▶  $a^{(t)} = Uh^{(t)} + b_2$
- ▶  $\hat{y}^{(t)} = \text{softmax}(a^{(t)})$
- ▶ Matrix dimensions get weird:

$$\frac{\partial a^{(t)}}{\partial U} \in \mathbb{R}^{k \times (k \times D_h)}$$

- ▶ But we don't need fancy tensors:

$$\nabla_U J^{(t)} = \left( \frac{\partial J^{(t)}}{\partial a^{(t)}} \cdot \frac{\partial a^{(t)}}{\partial U} \right)^T = \delta^{(2)(t)} (h^{(t)})^T \in \mathbb{R}^{k \times D_h}$$

- ▶ NumPy: `self.grads.U += outer(d2, hs[t])`

# Making sense of the madness

- ▶ Chain rule to the rescue!
- ▶  $a^{(t)} = Uh^{(t)} + b_2$
- ▶  $\hat{y}^{(t)} = \text{softmax}(a^{(t)})$
- ▶ Matrix dimensions get weird:

$$\frac{\partial a^{(t)}}{\partial U} \in \mathbb{R}^{k \times (k \times D_h)}$$

- ▶ But we don't need fancy tensors:

$$\nabla_U J^{(t)} = \left( \frac{\partial J^{(t)}}{\partial a^{(t)}} \cdot \frac{\partial a^{(t)}}{\partial U} \right)^T = \delta^{(2)(t)} (h^{(t)})^T \in \mathbb{R}^{k \times D_h}$$

- ▶ NumPy: `self.grads.U += outer(d2, hs[t])`



# Going deeper

- ▶ Really just need one simple pattern:
- ▶  $z^{(t)} = Hh^{(t-1)} + Wx^{(t)} + b_1$
- ▶  $h^{(t)} = f(z^{(t)})$
- ▶ Compute error delta ( $s = 0, 1, 2, \dots$ ):
  - ▶ From top:  $\delta^{(t)} = [h^{(t)} \circ (1 - h^{(t)})] \circ U^T \delta^{(2)(t)}$
  - ▶ Deeper:  $\delta^{(t-s)} = [h^{(t-s)} \circ (1 - h^{(t-s)})] \circ H^T \delta^{(t-s+1)}$
- ▶ These are just chain-rule expansions!

$$\frac{\partial J^{(t)}}{\partial z^{(t)}} = \frac{\partial J^{(t)}}{\partial a^{(t)}} \cdot \frac{\partial a^{(t)}}{\partial h^{(t)}} \cdot \frac{\partial h^{(t)}}{\partial z^{(t)}} = (\delta^{(t)})^T$$

# Going deeper

- ▶ Really just need one simple pattern:
- ▶  $z^{(t)} = Hh^{(t-1)} + Wx^{(t)} + b_1$
- ▶  $h^{(t)} = f(z^{(t)})$
- ▶ Compute error delta ( $s = 0, 1, 2, \dots$ ):
  - ▶ From top:  $\delta^{(t)} = [h^{(t)} \circ (1 - h^{(t)})] \circ U^T \delta^{(2)(t)}$
  - ▶ Deeper:  $\delta^{(t-s)} = [h^{(t-s)} \circ (1 - h^{(t-s)})] \circ H^T \delta^{(t-s+1)}$
- ▶ These are just chain-rule expansions!

$$\frac{\partial J^{(t)}}{\partial z^{(t)}} = \frac{\partial J^{(t)}}{\partial a^{(t)}} \cdot \frac{\partial a^{(t)}}{\partial h^{(t)}} \cdot \frac{\partial h^{(t)}}{\partial z^{(t)}} = (\delta^{(t)})^T$$

# Going deeper

- These are just chain-rule expansions!

$$\left. \frac{\partial J^{(t)}}{\partial b_1} \right|_{(t)} = \left( \frac{\partial J^{(t)}}{\partial a^{(t)}} \cdot \frac{\partial a^{(t)}}{\partial h^{(t)}} \cdot \frac{\partial h^{(t)}}{\partial z^{(t)}} \right) \cdot \frac{\partial z^{(t)}}{\partial b_1} = (\delta^{(t)})^T \frac{\partial z^{(t)}}{\partial b_1}$$

$$\left. \frac{\partial J^{(t)}}{\partial H} \right|_{(t)} = \left( \frac{\partial J^{(t)}}{\partial a^{(t)}} \cdot \frac{\partial a^{(t)}}{\partial h^{(t)}} \cdot \frac{\partial h^{(t)}}{\partial z^{(t)}} \right) \cdot \frac{\partial z^{(t)}}{\partial H} = (\delta^{(t)})^T \frac{\partial z^{(t)}}{\partial H}$$

$$\frac{\partial J^{(t)}}{\partial z^{(t-1)}} = \left( \frac{\partial J^{(t)}}{\partial a^{(t)}} \cdot \frac{\partial a^{(t)}}{\partial h^{(t)}} \cdot \frac{\partial h^{(t)}}{\partial z^{(t)}} \right) \cdot \frac{\partial z^{(t)}}{\partial h^{(t-1)}} = (\delta^{(t)})^T \frac{\partial z^{(t)}}{\partial z^{(t-1)}}$$

# Going deeper

- And there's shortcuts for them too:

$$\left( \frac{\partial J^{(t)}}{\partial b_1} \Big|_{(t)} \right)^T = \delta^{(t)}$$

$$\left( \frac{\partial J^{(t)}}{\partial H} \Big|_{(t)} \right)^T = \delta^{(t)} \cdot (h^{(t-1)})^T$$

$$\left( \frac{\partial J^{(t)}}{\partial z^{(t-1)}} \right)^T = \left[ h^{(t-1)} \circ (1 - h^{(t-1)}) \right] \circ H^T \delta^{(t)} = \delta^{(t-1)}$$

# Outline

## Backpropagation (continued)

- RNN Structure

- RNN Backpropagation


## Backprop on a DAG

- Example: Gated Recurrent Units (GRUs)

- GRU Backpropagation

# Motivation


- ▶ Gated units with “reset” and “output” gates
- ▶ Reduce problems with vanishing gradients



**Figure :** *You are likely to be eaten by a GRU.* (Figure from Chung, et al. 2014)

# Intuition

- ▶ Gates  $z_i$  and  $r_i$  for *each* hidden layer neuron
- ▶  $z_i, r_i \in [0, 1]$
- ▶  $\tilde{h}$  as “candidate” hidden layer
- ▶  $\tilde{h}, z, r$  all depend on  $x^{(t)}, h^{(t-1)}$
- ▶  $h^{(t)}$  depends on  $h^{(t-1)}$  mixed with  $\tilde{h}^{(t)}$



**Figure :** *You are likely to be eaten by a GRU.* (Figure from Chung, et al. 2014)

# Equations

- ▶  $z^{(t)} = \sigma(W_z x^{(t)} + U_z h^{(t-1)})$
- ▶  $r^{(t)} = \sigma(W_r x^{(t)} + U_r h^{(t-1)})$
- ▶  $\tilde{h}^{(t)} = \tanh(W x^{(t)} + r^{(t)} \circ U h^{(t-1)})$
- ▶  $h^{(t)} = z^{(t)} \circ h^{(t-1)} + (1 - z^{(t)}) \circ \tilde{h}^{(t)}$
- ▶ Optionally can have biases; omitted for clarity.




Figure : You are likely to be eaten by a GRU. (Figure from Chung, et al. 2014)

Same eqs. as Lecture 8, subscripts/superscripts as in Assignment #2.



# Backpropagation

Multi-path to compute  $\frac{\partial J}{\partial x^{(t)}}$

- ▶ Start with  $\delta^{(t)} = \left( \frac{\partial J}{\partial h^{(t)}} \right)^T \in \mathbb{R}^d$
- ▶  $h^{(t)} = z^{(t)} \circ h^{(t-1)} + (1 - z^{(t)}) \circ \tilde{h}^{(t)}$
- ▶ Expand chain rule into sum (*a.k.a. product rule*):

$$\begin{aligned} \frac{\partial J}{\partial x^{(t)}} &= \frac{\partial J}{\partial h^{(t)}} \cdot \left[ z^{(t)} \circ \frac{\partial h^{(t-1)}}{\partial x^{(t)}} + \frac{\partial z^{(t)}}{\partial x^{(t)}} \circ h^{(t-1)} \right] \\ &+ \frac{\partial J}{\partial h^{(t)}} \cdot \left[ (1 - z^{(t)}) \circ \frac{\partial \tilde{h}^{(t)}}{\partial x^{(t)}} + \frac{\partial (1 - z^{(t)})}{\partial x^{(t)}} \circ \tilde{h}^{(t)} \right] \end{aligned}$$

## It gets (a little) better

Multi-path to compute  $\frac{\partial J}{\partial x^{(t)}}$

- Drop terms that don't depend on  $x^{(t)}$ :

$$\begin{aligned}\frac{\partial J}{\partial x^{(t)}} &= \frac{\partial J}{\partial h^{(t)}} \cdot \left[ z^{(t)} \circ \frac{\partial h^{(t-1)}}{\partial x^{(t)}} + \frac{\partial z^{(t)}}{\partial x^{(t)}} \circ h^{(t-1)} \right] \\ &\quad + \frac{\partial J}{\partial h^{(t)}} \cdot \left[ (1 - z^{(t)}) \circ \frac{\partial \tilde{h}^{(t)}}{\partial x^{(t)}} + \frac{\partial (1 - z^{(t)})}{\partial x^{(t)}} \circ \tilde{h}^{(t)} \right] \\ &= \frac{\partial J}{\partial h^{(t)}} \cdot \left[ \frac{\partial z^{(t)}}{\partial x^{(t)}} \circ h^{(t-1)} + (1 - z^{(t)}) \circ \frac{\partial \tilde{h}^{(t)}}{\partial x^{(t)}} \right] \\ &\quad - \frac{\partial J}{\partial h^{(t)}} \frac{\partial z^{(t)}}{\partial x^{(t)}} \circ \tilde{h}^{(t)}\end{aligned}$$

# Almost there!

Multi-path to compute  $\frac{\partial J}{\partial x^{(t)}}$

- ▶ Now we really just need to compute two things:
- ▶ Output gate:

$$\frac{\partial z^{(t)}}{\partial x^{(t)}} = z^{(t)} \circ (1 - z^{(t)}) \circ W_z$$

- ▶ Candidate  $\tilde{h}$ :

$$\begin{aligned} \frac{\partial \tilde{h}^{(t)}}{\partial x^{(t)}} = & (1 - (\tilde{h}^{(t)})^2) \circ W \\ & + (1 - (\tilde{h}^{(t)})^2) \circ \frac{\partial r^{(t)}}{\partial x^{(t)}} \circ U h^{(t-1)} \end{aligned}$$

- ▶ Ok, I lied - there's a third.
- ▶ Don't forget to check all paths!

# Almost there!

Multi-path to compute  $\frac{\partial J}{\partial x^{(t)}}$

- ▶ Now we really just need to compute two things:
- ▶ Output gate:

$$\frac{\partial z^{(t)}}{\partial x^{(t)}} = z^{(t)} \circ (1 - z^{(t)}) \circ W_z$$

- ▶ Candidate  $\tilde{h}$ :

$$\begin{aligned} \frac{\partial \tilde{h}^{(t)}}{\partial x^{(t)}} = & \quad (1 - (\tilde{h}^{(t)})^2) \circ W \\ & + (1 - (\tilde{h}^{(t)})^2) \circ \frac{\partial r^{(t)}}{\partial x^{(t)}} \circ U h^{(t-1)} \end{aligned}$$

- ▶ Ok, I lied - there's a third.
- ▶ Don't forget to check all paths!

# Almost there!

Multi-path to compute  $\frac{\partial J}{\partial x^{(t)}}$

- ▶ Now we really just need to compute two things:
- ▶ Output gate:

$$\frac{\partial z^{(t)}}{\partial x^{(t)}} = z^{(t)} \circ (1 - z^{(t)}) \circ W_z$$

- ▶ Candidate  $\tilde{h}$ :

$$\begin{aligned} \frac{\partial \tilde{h}^{(t)}}{\partial x^{(t)}} = & \quad (1 - (\tilde{h}^{(t)})^2) \circ W \\ & + (1 - (\tilde{h}^{(t)})^2) \circ \frac{\partial r^{(t)}}{\partial x^{(t)}} \circ U h^{(t-1)} \end{aligned}$$

- ▶ Ok, I lied - there's a third.
- ▶ Don't forget to check all paths!

# Almost there!

Multi-path to compute  $\frac{\partial J}{\partial x^{(t)}}$


- ▶ Last one:

$$\frac{\partial r^{(t)}}{\partial x^{(t)}} = r^{(t)} \circ (1 - r^{(t)}) \circ W_r$$

- ▶ Now we can just add things up!
- ▶ (I'll spare you the pain...)


# Whew.

- ▶ Why three derivatives?
- ▶ Three arrows from  $x^{(t)}$  to distinct nodes
- ▶ Four paths total ( $\frac{\partial z^{(t)}}{\partial x^{(t)}}$  appears twice)



# Whew.

- ▶ GRUs are complicated
- ▶ All the pieces are simple
- ▶ Same matrix gradients that you've seen before





# Summary

- ▶ Check your dimensions!
- ▶ Write error vectors  $\delta$ ; just parentheses around chain rule
- ▶ Combine simple operations to make complex network
  - ▶ Matrix-vector product
  - ▶ Activation functions (tanh, sigmoid, softmax)

Good luck on Wednesday!



May the 4th be with you!

# CS 224d Midterm Review (word vectors)

Peng Qi

May 5, 2015

# Outline


- ▶ word2vec and GloVe revisited
- ▶ word2vec with backpropagation

# Outline

- ▶ word2vec and GloVe revisited
  - ▶ Skip-gram revisited
  - ▶ (Optional) CBOW and its connection to Skip-gram
  - ▶ (Optional) word2vec as matrix factorization (conceptually)
  - ▶ GloVe v.s. word2vec
- ▶ word2vec with backpropagation

# Skip-gram

- ▶ *Task*: given a **center word**, predict its **context words**
- ▶ For each word, we have an “**input vector**”  $v_w$  and an “**output vector**”  $v'_w$



# Skip-gram


- ▶ We have seen two types of costs for an expected word given a vector prediction  $r$

$$CE(w_i|r) = -\log \left( \frac{\exp(r^\top v'_{w_i})}{\sum_{j=1}^{|V|} \exp(r^\top v'_{w_j})} \right)$$

$$\begin{aligned} NEG(w_i|r) = & -\log(\sigma(r^\top v'_{w_i})) \\ & - \sum_{k=1}^K \log(\sigma(-r^\top v'_{w_k})) \end{aligned}$$

In the case of skip-gram, the vector prediction  $r$  is just the “input vector” of the center word,  $v_{w_i}$ .

$\sigma(\cdot)$  is the sigmoid (logistic) function.




# Skip-gram

- Now we have all the pieces of skip-gram, the cost for a context window  $[w_{i-C}, \dots, w_{i-1}, w_i, w_{i+1}, \dots, w_{i+C}]$  is ( $w_i$  is the center word)

$$J_{\text{skip-gram}}([w_{i-C}, \dots, w_{i+C}]) = \sum_{i-C \leq j \leq i+C, i \neq j} F(w_j | v_{w_i})$$

where  $F$  is one of the cost functions we defined in the previous slide.


- You might ask: but why are we introducing so many notations?





# Skip-gram v.s. CBOW

## Skip-gram




Task

Center word  $\rightarrow$  Context

$r$

$v_{w_i}$

## CBOW



Context  $\rightarrow$  Center word

$f(v_{w_{i-C}}, \dots, v_{w_{i-1}}, v_{w_{i+1}}, \dots, v_{w_{i+C}})$

## word2vec as matrix factorization (conceptually)

- ▶ Matrix factorization

$$\begin{bmatrix} M \end{bmatrix}_{n \times n} \approx \begin{bmatrix} A^\top \end{bmatrix}_{n \times k} \begin{bmatrix} B \end{bmatrix}_{k \times n}$$
$$M_{ij} \approx a_i^\top b_j$$

- ▶ Imagine  $M$  is a matrix of counts for events co-occurring, but we only get to observe the co-occurrences one at a time. E.g.

$$M = \begin{bmatrix} 1 & 0 & 4 \\ 0 & 0 & 2 \\ 1 & 3 & 0 \end{bmatrix}$$

but we only see

$(1,1), (2,3), (3,2), (2,3), (1,3), \dots$

## word2vec as matrix factorization (conceptually)

$$M_{ij} \approx a_i^\top b_j$$

- ▶ Whenever we see a pair  $(i, j)$  co-occur, we try to increasing  $a_i^\top b_j$
- ▶ We also try to make all the other inner-products smaller to account for pairs never observed (or unobserved yet), by decreasing  $a_{-i}^\top b_j$  and  $a_i^\top b_{-j}$
- ▶ Remember from the lecture that the word co-occurrence matrix usually captures the semantic meaning of a word?  
For word2vec models, roughly speaking,  $M$  is the windowed word co-occurrence matrix,  $A$  is the output vector matrix, and  $B$  is the input vector matrix.
- ▶ Why not just use one set of vectors? It's equivalent to  $A = B$  in our formulation here, but less constraints is usually easier for optimization.

## GloVe v.s. word2vec

	Fast training	Efficient usage of statistics	Quality affected by size of corpora	Captures complex patterns
Direct prediction (word2vec)	Scales with size of corpus	No	No*	Yes
GloVe	Yes	Yes	No	Yes

\* Skip-gram and CBOW are qualitatively different when it comes to smaller corpora

# Outline

- ▶ word2vec and GloVe revisited
- ▶ word2vec with backpropagation

## word2vec with backpropagation



$$CE(w_i|r) = -\log \left( \frac{\exp(r^\top v'_{w_i})}{\sum_{j=1}^{|V|} \exp(r^\top v'_{w_j})} \right)$$



$$\begin{aligned} CE(w_i|r) &= CE(\hat{y}, y_i) \\ \hat{y} &= \text{softmax}(\theta) \\ \theta &= (V')^\top r \end{aligned}$$



$$\delta = \frac{\partial CE}{\partial \theta} = \hat{y} - y_i$$



$$\begin{aligned} \frac{\partial CE}{\partial V'} &= r \delta^\top \\ \frac{\partial CE}{\partial r} &= V' \delta \end{aligned}$$

Thanks for your attention  
and best of luck with the mid-term!



Any questions?